

Purpose of Read Mapping

- Identify the origin of a read
- RNA-seq: Infer the gene that was expressed and that generated the read
- BUT: Is it always possible to find the read origin?

Sequence Alignment

- Optimization problem: Find the alignment with the **highest score**
- Global alignment (end-to-end): Needleman-Wunsch
- Local alignment: Smith-Waterman

Score matrix:

	a	c	g	t
a	+2	-2	-1	-2
c	-2	+2	-2	-1
g	-1	-2	+2	-2
t	-2	-1	-2	+2

Gap score: -3

c	c	t	a	g	t	t	c	a	t
				x		x			
t	g	t	a	a	t	-	c	a	c
		+2	+2	-1	+2	-3	+2	+2	

Alignment score = 6


Local alignments (Smith-Waterman)

- Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)
- Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system
- Need to set scores for **match**, penalties for **mismatch** and **gap** (typically, mismatch penalties are set according to evolutionary knowledge and error profiles)

- Match: +1; Mismatch: -1; Gap: -2

Reference Sequence

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0					
T	0					
G	0					
A	0					
C	0					



$$H(i, j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion

Local alignments (Smith-Waterman)

- Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)
 - Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system
 - Need to set scores for **match**, penalties for **mismatch** and **gap**
-
- Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1			
T	0					
G	0					
A	0					
C	0					

$$H(i,j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1,j-1) + w(a_i,b_j) \\ H(i-1,j) + w(a_i,-) \\ H(i,j-1) + w(-,b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion

Local alignments (Smith-Waterman)

- Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)
 - Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system
 - Need to set scores for **match**, penalties for **mismatch** and **gap**
-
- Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1	0	0	0
T	0	0	0	2	0	1
G	0					
A	0					
C	0					

$$H(i,j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1,j-1) + w(a_i,b_j) \\ H(i-1,j) + w(a_i,-) \\ H(i,j-1) + w(-,b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion

Local alignments (Smith-Waterman)

- Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)
- Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system
- Need to set scores for **match**, penalties for **mismatch** and **gap**

- Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1	0	0	0
T	0	0	0	2	0	1
G	0	0	0	0	3	1
A	0	1	1	0	1	2
C	0	0	0	0	0	0

$$H(i, j) = \max \begin{cases} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{cases}, 1 \leq i \leq m, 1 \leq j \leq n$$


Match/Mismatch

Deletion

Insertion

Local alignments (Smith-Waterman)

- Traceback:
- Start with maximum score (here, 3)
- Follow path that gives multiple score



$$H(i,j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1,j-1) + w(a_i,b_j) \\ H(i-1,j) + w(a_i,-) \\ H(i,j-1) + w(-,b_j) \end{array} \right\}, 1 \leq i \leq m, 1 \leq j \leq n$$


Match/Mismatch

Deletion

Insertion

Local alignments (Smith-Waterman)

- Traceback:
- Start with maximum score (here, 3)
- Follow path that gives multiple score
- This gives:
- AATGT-
- ATGAC



$$H(i,j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1,j-1) + w(a_i,b_j) \\ H(i-1,j) + w(a_i,-) \\ H(i,j-1) + w(-,b_j) \end{array} \right\}, 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion


Local alignments (Smith-Waterman)

- Other considerations:
 - Natural extension is to have a different *gap opening* and a *gap extension* penalty (former generally being larger)
 - GO penalty=-2; GE penalty=-1; Mismatch=-1; Match=1
 - With above penalties, which is the best scoring alignment?
- | | | |
|-----------|---------|----------|
| • AT-C-GT | ATC--GT | AT-C--GT |
| • ATTTTGT | ATTTTGT | ATT-TTGT |

- Dynamic programming is slow
- Speed-up with heuristics
 - e.g. exactly align short subsequences and extend these alignments
 - e.g. BLAST / BLAT
 - no longer guaranteed to find the best alignments
 - exact matches are found by index lookup


Hashing based alignment

a. Seed extraction from reference genome




indexing

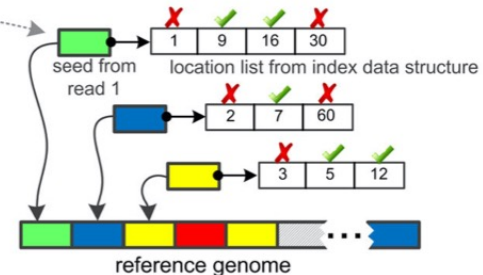
b. Seed indexing using suffix tree or hash table




c. Seed extraction from reads



d. Seed querying and filtering



e. Seed chaining and pre-alignment filtering



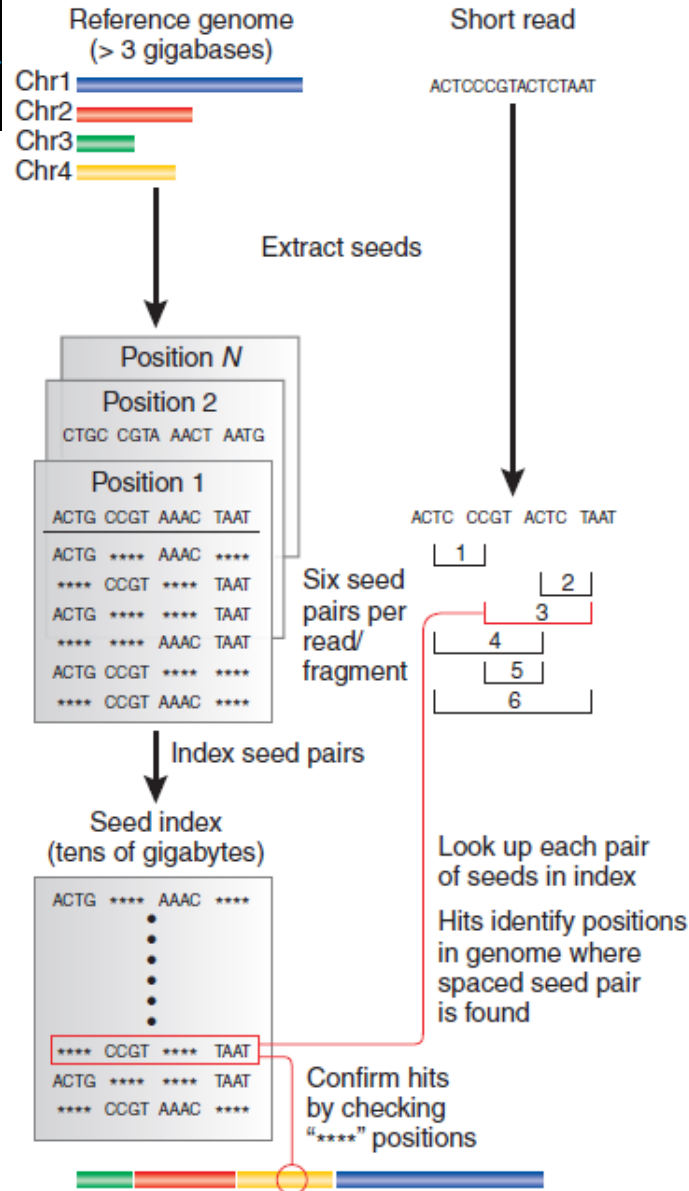
f. Alignment verification

	C	G	T	T	A	G	T	C	T	A	...
C	0	0	0	0	0	0	0	0	0	0	...
G	0	2	2	2	2	2	2	2	2	2	...
T	0	2	3	3	3	3	3	4	4	4	...
A	0	2	3	5	5	5	5	5	6	6	...
T	0	2	3	5	7	7	7	7	7	7	...
A	0	3	3	5	7	9	9	9	9	9	...
G	0	2	4	5	7	9	11	11	11	11	...
T	0	2	4	6	7	9	11	13	13	13	...
A	0	2	4	6	7	9	11	13	14	14	...
T	0	2	4	6	8	9	11	13	14	16	...

.bam/.sam file contains necessary alignment information (e.g., type, location, and number of each edit)


Index Genome: Spaced Seeds

- Tags and tag-sized pieces of reference are cut into small “seeds.”
- Pairs of spaced seeds are stored in an index.
- Look up spaced seeds for each tag.
- For each “hit,” confirm the remaining positions.



Index Genome: Burrows-Wheeler Transform

- Store entire reference genome.
- Align tag base by base from the end.
- When tag is traversed, all active locations are reported.
- If no match is found, then back up and try a substitution.



[https://www.geeksfor
geeks.org/inverting-
burrows-wheeler-](https://www.geeksforgeeks.org/inverting-burrows-wheeler-)

The “Last-First mapping” property

The relative ordering of a particular character (say **c**) in column 1 is the same as that in the last column

c_1 a c_2 a a c_3 g \$

\$	c_1	a	c_2	a	a	c_3	g
a	a	c_3	g	\$	c_1	a	c_2
a	c_2	a	a	c_3	g	\$	c_1
a	c_3	g	\$	c_1	a	c_2	a
c_2	a	a	c_3	g	\$	c_1	a
c_1	a	c_2	a	a	c_3	g	\$
c_3	g	\$	c_1	a	c_2	a	a
g	\$	c_1	a	c_2	a	a	c_3

The “Last-First mapping” property

The relative ordering of a particular character (say **c**) in column 1 is the same as that in the last column

c_1 a c_2 a a c_3 g \$

\$	c_1	a	c_2	a	a	c_3	g
a	a	c_3	g	\$	c_1	a	c_2
a	c_2	a	c_3	g	\$	c_1	c_2
a	c_3	g	\$	c_1	a	c_2	a
c_2	a	a	c_3	g	\$	c_1	a
c_1	a	c_2	a	a	c_3	g	\$
c_3	g	\$	c_1	a	c_2	a	a
g	\$	c_1	a	c_2	a	a	c_3

Proof:

Suppose cX and cY are cyclic permutations of the input T .

Suppose $cX < cY$ (in lexicographical ordering)

Then $Xc < Yc$ (in lexicographical ordering)

The LF-mapping property follows.

BWT is reversible

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c



BWT is reversible

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

g\$ ca ca aa ac \$c ac cg

➔

\$c aa ac ac ca ca cg g\$

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

- Range ← range of last character in 1st column
- While characters left (and nonzero range):
 - Lookup first and last match to preceding character in final column
 - Range ← LF-mapping of first and last match

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c


Range ← range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range ← LF-mapping of first and last match

Lookup AAC



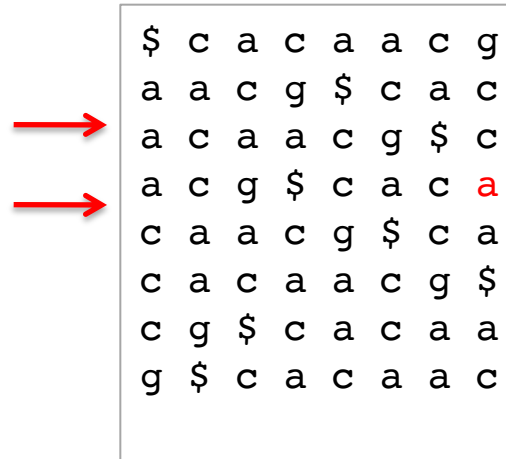
Range ← range of last character in 1st column
 While characters left (and nonzero range):
 Lookup first and last match to preceding character in final column
 Range ← LF-mapping of first and last match

Lookup AAC

	\$	c	a	c	a	a	c	g
	a	a	c	g	\$	c	a	c
→	a	c	a	a	c	g	\$	c
→	a	c	g	\$	c	a	c	a
	c	a	a	c	g	\$	c	a
	c	a	c	a	a	c	g	\$
	c	g	\$	c	a	c	a	a
	g	\$	c	a	c	a	a	c

Range ← range of last character in 1st column
While characters left (and nonzero range):
 Lookup first and last match to preceding character in final column
 Range ← LF-mapping of first and last match

Lookup AAC



\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c


Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup AAC



Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup AAC

→

→

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Range ← range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range ← LF-mapping of first and last match

Comparison

Spaced seeds

- Requires ~50Gb of memory.
- More straightforward to program.
- Example:
 - Subread
- More tolerant to
 - sequence variations
 - sequencing errors

Burrows-Wheeler

- Compressed index requires <2Gb of memory for human genome.
- can be very fast
- More complicated to program.
- Examples:
 - bowtie2
 - BWA

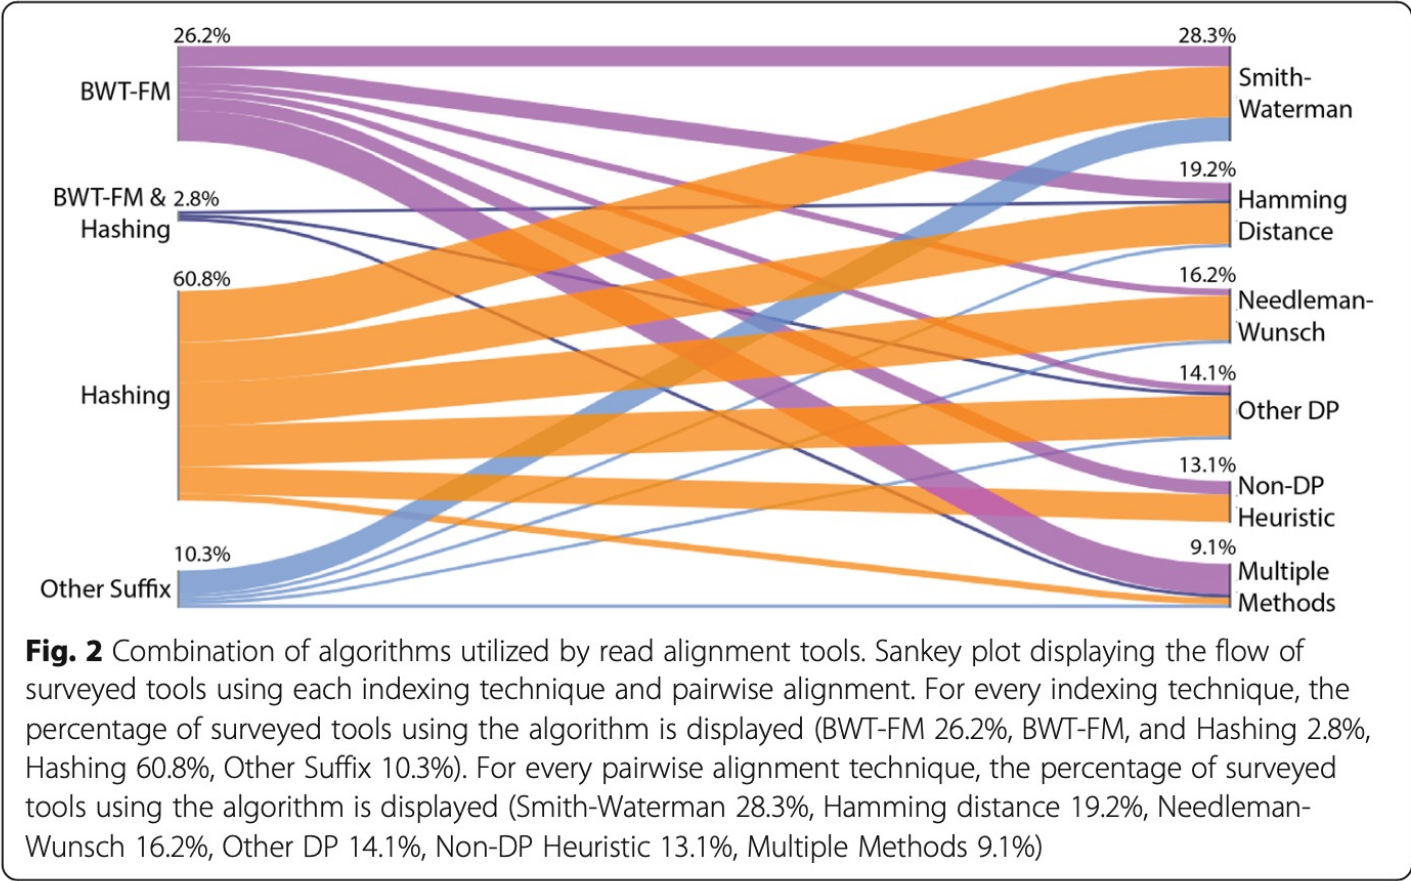
Characteristics of algorithms

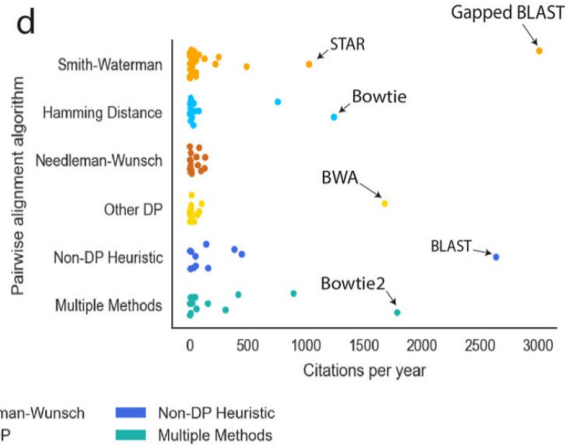
	Hashing	Suffix tree and BWT-FM
Easy to implement	Yes	No
Search for exact/inexact match	Exact	Exact and inexact
Index size	Large	Compressed (small)
Indexing time	Small	Large
Seed query speed	O(1), fast	Slow
Seed length	Fixed length per index	Can be fixed or variable

Alser et al.

<https://genomebiology.biomedcentral.com/articles/10.1186/s13059-021-02443-7>

Combining index lookup and pairwise alignment





Alignment with Mismatches

- Mismatches can occur because of
 - sequencing error (error rate $\sim 1/500$)
 - mutation; (human mutation rate $\sim 1/1e4$)
 - \rightarrow if reads are long ($> 100nt$) reads with mismatches will not be rare; more than $\sim 10\%$ of the reads may have a mismatch
- If there is a sequence mismatch the index lookup fails! How to find nevertheless an alignment?

- Strategies:
 - If the BWT lookup fails at position k , try all different bases at position k
→ drawback: computing effort grows exponentially with the number of mismatches
→ implemented e.g. in bowtie
 - Chop reads in segments (seeds) and align those mismatch-free and stitch seed alignments together
→ implemented e.g. in bowtie

Consideration: Sequencing Errors – PHRED quality

- Each called base is given a quality score Q
- $Q = 10 * \log_{10}(\text{estimated probability that call is wrong})$

10 prob = 0.1

20 prob = 0.01

30 prob = 0.001

[Q30 often used as a threshold for useful sequence data]


→ down-weight low-quality bases when computing the alignment score

- These quality scores are called PHRED scores.
- PHRED scores are determined by the sequencer that directly rates how reliable the measured signal is

Considerations: Read trimming

- Sequencers have systematic errors
- Illumina sequencers have a higher error rate at the first few bases (1-5)
- Basically, all sequencers have increasing error rate towards the end of the read
- Hard trimming: trim a fixed number of bases from the beginning and/or end
- Quality trimming: cut the end of the read as soon as the base quality drops below a threshold

Considerations: Read trimming



- If the inserted DNA/RNA fragment is too short, the read will contain part of the adapter
- Adapter trimming can be challenging if
 - if the insert is 90 – 100 bases
 - if the read has many sequencing errors

Multiple Alignments

- A read may have multiple valid alignments with identical or similarly good alignment scores
- Aligners allow to choose different reporting strategies:
 - Randomly select one alignment from the top-scoring alignments
 - Report all alignments that are within *delta* of the top-scoring alignment; clip if more than *Nmax* alignments are found
 - Report only alignments if they are unique (no other alignment within *delta* of the alignment score)
 - Do not report anything if more than *Nmax* valid alignments are found
 - ...
- Whether a read has a **unique alignment** depends on
 - the read sequence and the sequence homology of the organism
 - the search algorithm of the aligner
 - the reporting strategy of the aligner

Read Alignment Summary (I)

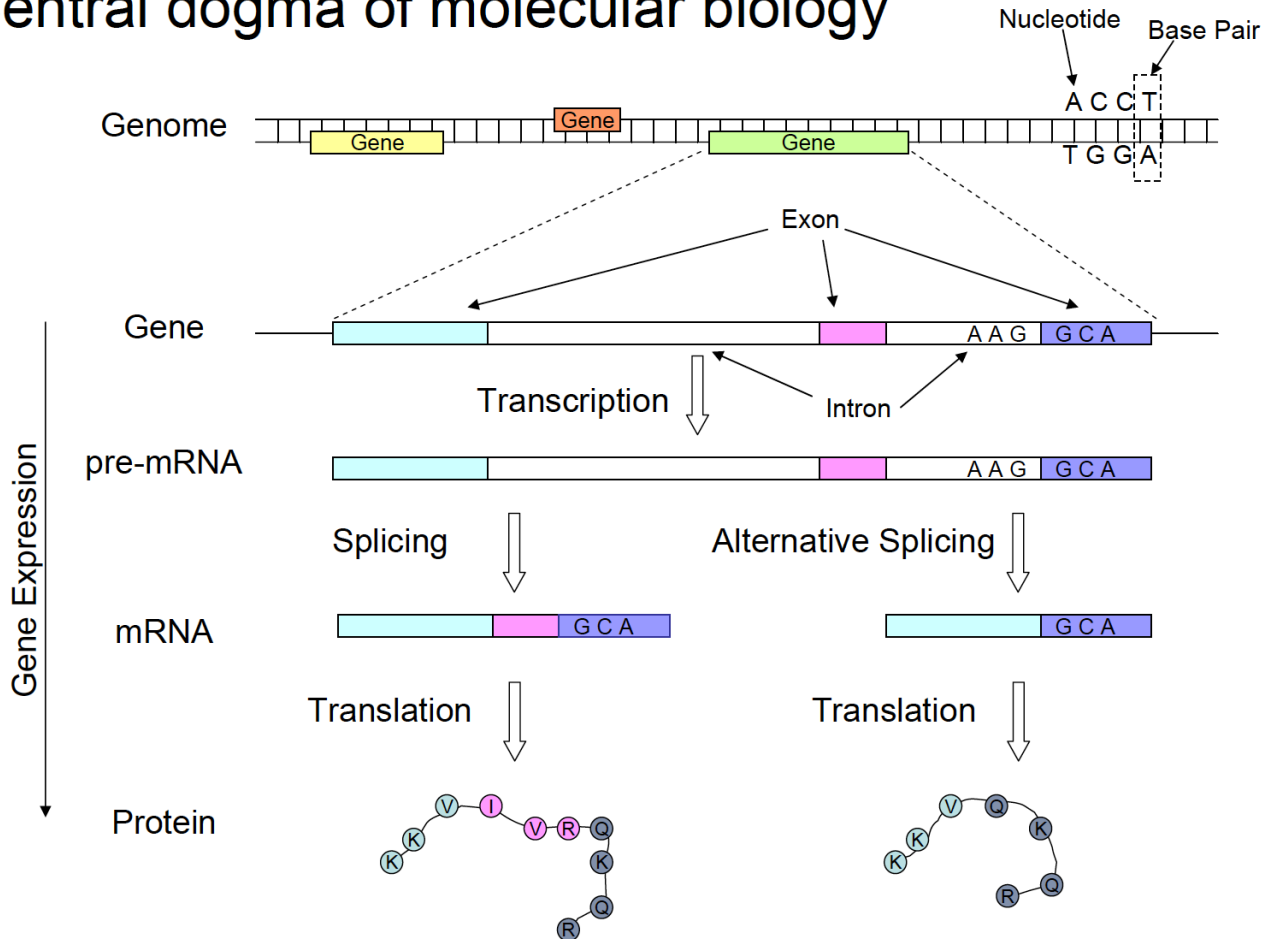
How to map billions of reads?

- The major aligners use the Burrows-Wheeler-Transform to create an index of the reference. Even for the human genome the index fits into 3GB RAM.
- Reads are aligned by **index lookup** not by sequence comparison
- The lookup of a perfect match read is faster than loading the read and writing the alignment coordinates to the output file!
- If the lookup fails because of SNPs or sequencing errors, an actual sequence comparison is done –and this can take time!

Read Alignment Summary (II)

- Alignment is a score optimization problem
- Aligners find the alignment with the **highest alignment score**
- Aligners do not run an extensive search, they use **index lookup and heuristics** to find the alignment with highest score; these heuristics are not guaranteed to find the alignment with the maximum score
- If the heuristics are well chosen and the scores are well defined, then the alignment will be very often the highest-scoring alignment and that will be the correct alignment

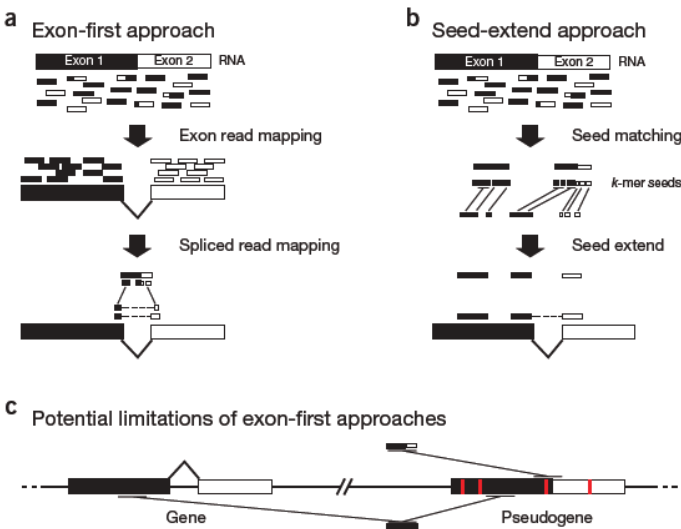
Central dogma of molecular biology



92–94% of human genes undergo alternative splicing,

RNA-seq Mapping

- Mapping targets:
 - transcriptome
 - splice junction library
 - genome
- Mouse retina 60+60bp reads:
 - 41 Mio of 91 Mio map to junctions

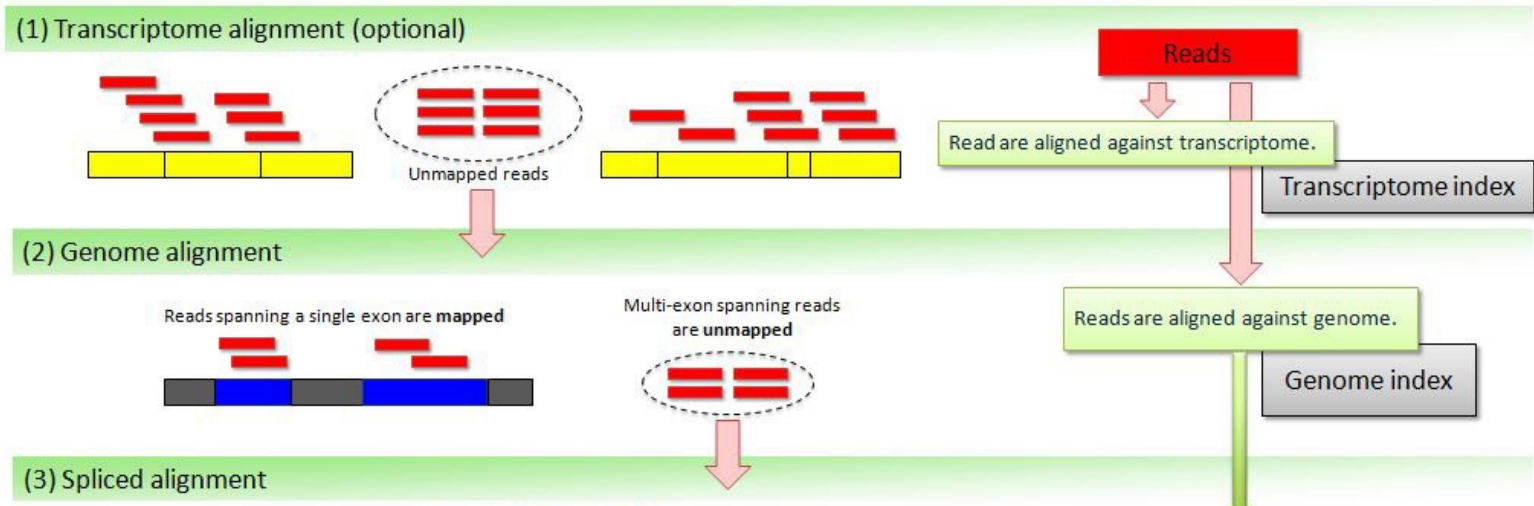


Strategies:

- Exon-first: fast
- Seed-extend:
 - good with polymorphisms
 - simultaneous spliced/unspliced mapping






Tophat2 pipeline – illustrative of the challenge of mapping RNA-seq reads

- Note: Tophat2 is outdated, use STAR, HISAT, Subread, ...



(3) Spliced alignment

(3-1) Segment alignment to genome

(3-2) Identification of splice sites
(including indels and fusion break points)(3-3) Segments aligned to junction
flanking sequences(3-4) Segment alignments stitched
together to form whole read alignments(3-5) Re-alignment of reads minimally
overlapping intronsReads are split into smaller segments
which are then aligned to the genome.

Genome index




Segment mappings are used to find potential splice sites
usually when the distance between the mapped positions
of the left and the right segments are longer than the
length of the middle part of a read.Sequences flanking a splice site are concatenated
and segments are aligned to them.


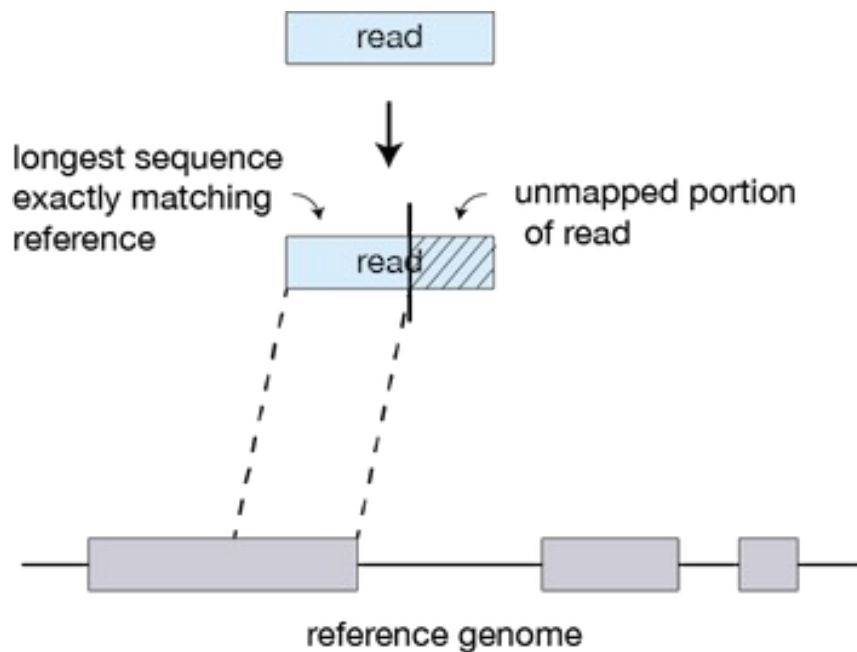
Junction flanking index



Mapped segments against either genome or flanking
sequences are gathered to produce whole read alignments.Genome mapped reads with alignments extending a few
bases into introns are re-aligned to exons instead.


STAR: universal RNA-seq aligner

- Designed to align the non-contiguous sequences directly to the reference genome
- Steps:
 - Search Maximal Mappable Prefix (MMP)
 - clustering/stitching/scoring









- Tradeoff:
 - speed
 - RAM
 - accuracy
 - sensitivity

- <http://www.ecseq.com/support/benchmark>

Performance comparison of RNA-seq mappers

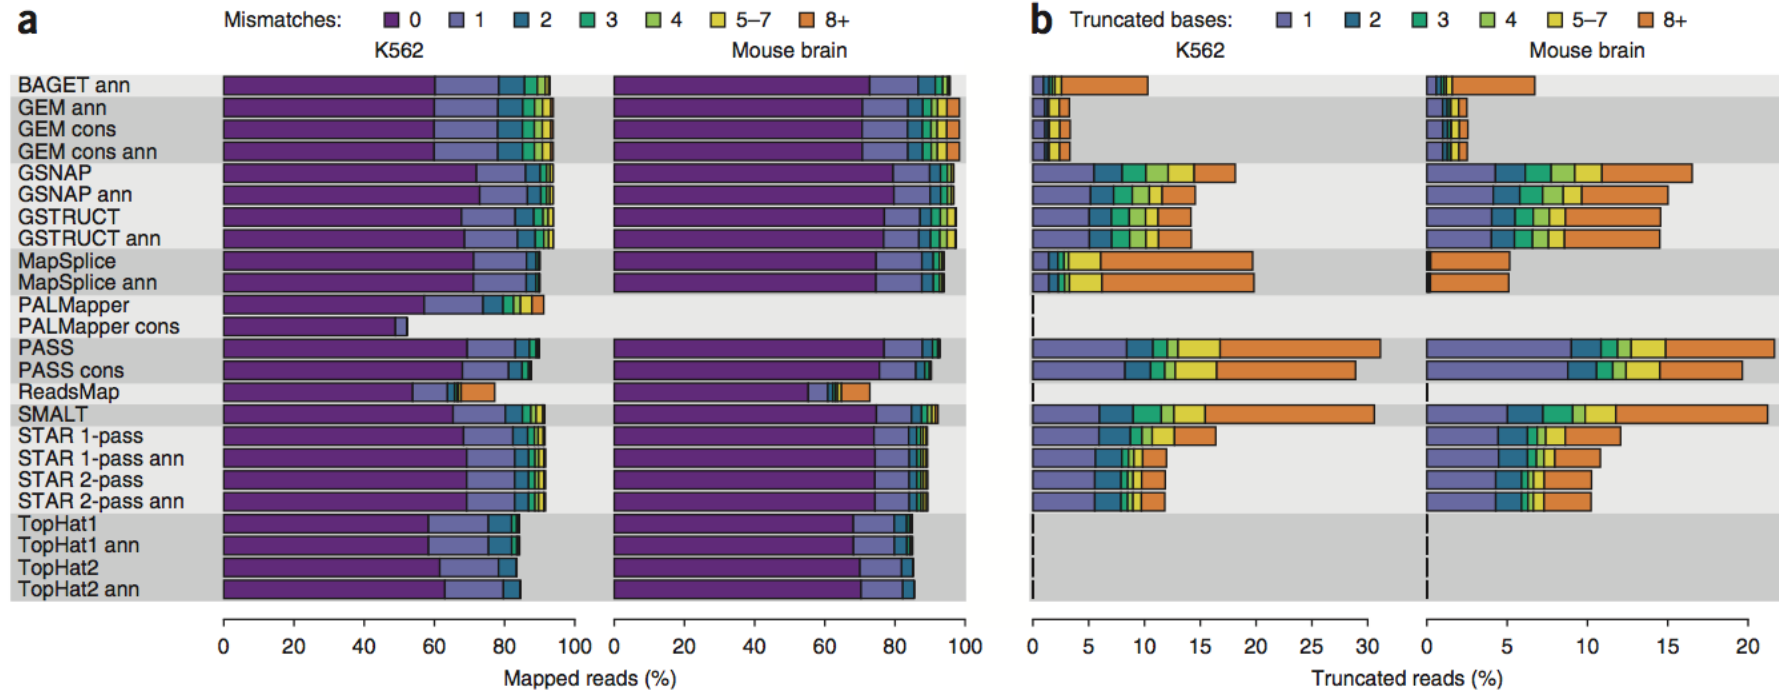


Figure 2 | Mismatch and truncation frequencies. **(a)** Percentage of sequenced reads mapped with the indicated number of mismatches. **(b)** Percentage of sequenced reads truncated at either or both ends. Bar colors indicate the number of bases removed.

Engström et al. Systematic evaluation of spliced alignment programs for RNA-seq data. Nat Methods. 2013.


STAR 2-pass mode

- goal: improve alignment around novel splice junctions
- does not discover more splice junctions but aligns more reads to discovered splice junctions
- initial mapping might fail to align reads where only a short part of the read spans the splice junction

How long should a read be so that it can be uniquely aligned??

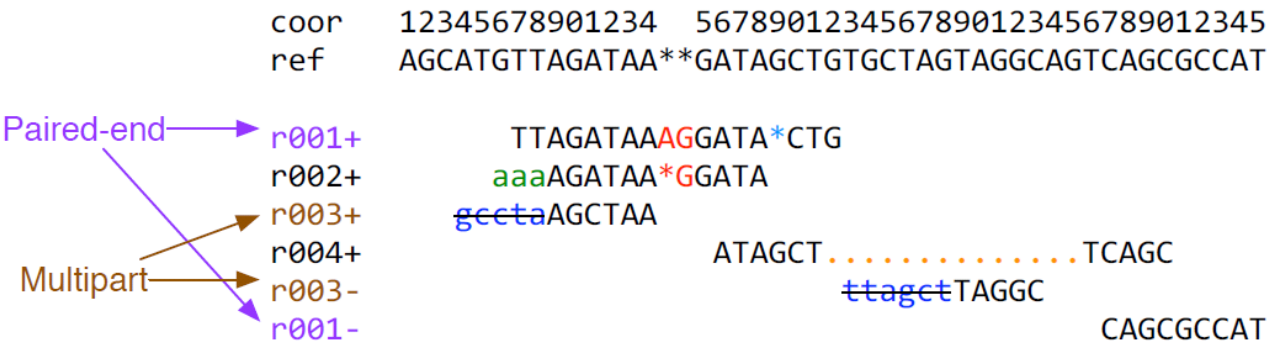
Current practice

- use ≥ 75 nt for aligning RNA-reads to mammals
- use $\geq 100+100$ nt for aligning DNA reads for variant calling
- use 21nt for aligning micro-RNA reads to miRbase



- **SAM (Sequence Alignment/Map)**
 - Single unified format for storing read alignments to a reference genome
 - Large plain text file
- **BAM (Binary Alignment/Map)**
 - Binary equivalent of SAM
 - Compressed data plus index (bai)
 - Developed for fast processing/indexing

An example of read mapping



Corresponding SAM file

Paired-end

Multipart

r001+

r002+

r003+

r004+

r003-

r001-

12345678901234

5678901234567890123456789012345

AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT

TTAGATAAAGGATA*CTG

aaaAGATAA*GGATA

gctaAGCTAA

ATAGCT.....TCAGC

ttagctTAGGC

CAGCGCCAT

Ins & padding

Soft clipping

Splicing

Hard clipping

@SQ SN:ref LN:45

r001 163 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTA *

r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *

r003 0 ref 9 30 5H6M * 0 0 AGCTAA * NM:i:1

r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC *

r003 16 ref 29 30 6H5M * 0 0 TAGGC * NM:i:0

r001 83 ref 37 30 9M = 7 -39 CAGCGCCAT *

QNAME

RNAME

MAPQ?

MRNM

ISIZE

QUAL

FLAG

POS

CIGAR?

MPOS

SEQ

Slide modified from Heng Li's presentation

CIGAR string - compact representation of an alignment

- M - match or mismatch
- I – insertion
- D – deletion
- S - soft clip
 - Clipped sequences stored in SAM
- H - hard clip
 - Clipped sequences not stored in SAM
- N – skipped reference bases, splicing

Match/mismatch, indels

Ref: ACGCA**T**G--GT
 Read: ATGCA-TG**C**AGT
 Cigar: 5M1D2M2I2M

Soft clipping

REF: ATCGTGTAACCTGACTAGTTAA
 READ: gggGTGTAACC-GACTAGggg
 Cigar: 3S8M1D5M4S

Hard clipping


REF: ATCGTGTAACCTGACTAGTTAA
 READ: gggGTGTAACC-GACTAGggg
 Cigar: 3H8M1D5M4H

Mapping quality – the MAPQ tag


- The MAPQ tag does indicate the quality of an alignment
- Equivalent to the PHRED score for base
- However – the error probability can not be as accurately determined as the base quality
- The different aligners have different methods for assigning the quality

Mapping QC: Mapping statistics

- How well did my sequence library align to my reference?
- Summary statistics (per read library)
 - % reads with unique alignment
 - % reads with multiple alignment
 - % reads with no alignment
 - % reads properly paired (for paired end libraries)




Mapping QC: Profiling efficiency by transcript annotation




- Transcript annotation: intron, exon, up/down stream, unannotated
- Expression profiling efficiency

Mapping QC: Coverage bias




Mapping QC: Junction saturation





- Depth needed for alternative splicing analysis
- All annotated splice junctions are detected - a saturated RNA-seq dataset

Statistics on Trimming, Soft Clipping and Mismatches



- Mismatches depend on position in the read
- If bases 1 or 2 do not match, the STAR aligner soft-clips them

Statistics on Trimming, Soft Clipping and Mismatches



- Example from a low quality sample (low input <1ng, slight degradation)
- Short inserts, increased mismatch rate

Where mappers have issues

- paralog genes that exist on X and Y chromosomes
- MHC genes (part of the adaptive immune response)
 - poly-morphic: several alleles (versions) exist on that gene's locus

RNA-seq mapping QC tools

- RNA-SeQC
 - <https://github.com/getzlab/rnaseqc>
- RSeQC
 - <http://rseqc.sourceforge.net/>